

# CSC 4780/6780

## Fall 2022

### Homework 10

October 25, 2022

This homework is due at 11:59 pm on Sunday, Nov 6. It must be uploaded to iCollege by then. No credit will be given for late submissions. A solution will be released by noon on Monday, Nov 7

it is always a good idea to get this done and turned in early. You can turn it in as many times as you like – iCollege will only keep the last submission. If, for some reason, you are unable to upload your solution, email it to me before the deadline.

Incidentally, I rarely check my iCollege mail, but I check my `dhillegass@gsu.edu` email all the time. Send messages there.

Be sure to rename your solution directory to match your name.

## 1 AutoML for Regression

I once worked with an old engineer who would quietly listen to younger engineers arguing over what each thought was the best solutions to a problem. Eventually, he would say, "There is no point in arguing about things that can be tested." And then he would go and do an experiment that ended the argument.

As we get better and better at working with these models, we can begin to guess which will be best. However, a lot of the time we can just try all of them.

In this exercise, you will get a data set for regression and you will use `pycaret` to find the best candidates and test them against each other.

### 1.1 Training and Comparing

`train_concrete.csv` and `test_concrete.csv` contain data about the compressive strength of several different concrete mixes: <https://archive.ics.uci.edu/ml/datasets/Concrete+Compressive+Strength>

You will write a program called `concrete_train.py` that will use `pycaret's compare_models` (no turbo!) to try a large variety of regression algorithms on `train_concrete.csv`.

It will pick the best six (based on  $R^2$ ) and it will tune (using at least 24 different parameter combinations) and finalize each before saving the finalized model to a pickle file. Thus six `.pkl` files will be created.

Run the program and save the output to `train.txt`.

`train.txt` should look like this:

```
*** Setting up session***
      Description      Value
0      Session id      8371
1      Target          csMPa
2      Target type      Regression
...
18     USI              6171
*** Set up: 1.89 seconds
```

	Model	MAE	MSE	RMSE	\
catboost	CatBoost Regressor	2.8945	18.6345	4.2833	
lightgbm	Light Gradient Boosting Machine	3.4278	24.0534	4.8647	
et	Extra Trees Regressor	3.5456	26.9685	5.1605	
rf	Random Forest Regressor	3.8756	27.8363	5.2477	
gbr	Gradient Boosting Regressor	3.9316	28.5122	5.3121	
mlp	MLP Regressor	5.1949	46.8561	6.8208	
dt	Decision Tree Regressor	5.0301	56.0932	7.3678	
ada	AdaBoost Regressor	6.3680	61.0570	7.7998	
knn	K Neighbors Regressor	7.3850	96.5281	9.7726	
br	Bayesian Ridge	8.1946	108.8475	10.4094	
kr	Kernel Ridge	8.2325	108.9195	10.4127	
en	Elastic Net	8.2139	109.0079	10.4165	
ridge	Ridge Regression	8.2163	109.0042	10.4161	
lr	Linear Regression	8.2163	109.0043	10.4161	
lasso	Lasso Regression	8.2147	109.0795	10.4200	
ard	Automatic Relevance Determination	8.2609	109.4030	10.4368	
huber	Huber Regressor	8.1080	116.0969	10.6962	
par	Passive Aggressive Regressor	9.6928	149.4162	12.0777	
lar	Least Angle Regression	9.9147	163.0199	12.6530	
omp	Orthogonal Matching Pursuit	12.0965	216.9526	14.6893	
svm	Support Vector Regression	12.0595	227.7054	15.0594	
tr	TheilSen Regressor	9.0357	232.6367	14.6477	
llar	Lasso Least Angle Regression	13.6897	286.5223	16.8957	
dummy	Dummy Regressor	13.6897	286.5223	16.8957	
ransac	Random Sample Consensus	10.4945	352.2832	17.8668	

  

	R2	RMSLE	MAPE	TT (Sec)
catboost	0.9337	0.1358	0.1003	0.227

lightgbm	0.9143	0.1576	0.1191	0.016
et	0.9043	0.1622	0.1235	0.032
rf	0.9010	0.1772	0.1404	0.035
gbr	0.8986	0.1760	0.1383	0.016
mlp	0.8327	0.2217	0.1769	0.090
dt	0.8006	0.2311	0.1713	0.007
ada	0.7840	0.2828	0.2631	0.017
knn	0.6541	0.3188	0.2843	0.007
br	0.6097	0.3320	0.3135	0.007
kr	0.6092	0.3307	0.3135	0.009
en	0.6091	0.3315	0.3134	0.090
ridge	0.6091	0.3312	0.3131	0.089
lr	0.6091	0.3312	0.3131	0.219
lasso	0.6088	0.3318	0.3137	0.095
ard	0.6073	0.3314	0.3149	0.007
huber	0.5809	0.3235	0.3038	0.010
par	0.4758	0.3902	0.3636	0.007
lar	0.4182	0.4281	0.3720	0.007
omp	0.2359	0.4757	0.5022	0.007
svm	0.1996	0.4822	0.5051	0.009
tr	0.1558	0.3313	0.3066	0.171
llar	-0.0068	0.5397	0.6003	0.007
dummy	-0.0068	0.5397	0.6003	0.006
ransac	-0.2651	0.3615	0.3362	0.017

\*\*\* compare\_models: 16.59 seconds

\*\*\* Best:

```

CatBoostRegressor
LGBMRegressor
ExtraTreesRegressor
RandomForestRegressor
GradientBoostingRegressor
MLPRegressor

```

\*\*\* 0 - CatBoostRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						
0	3.4779	28.4514	5.3340	0.9139	0.1772	0.1307
...						
9	2.5817	15.6806	3.9599	0.9387	0.1492	0.1040
Mean	3.0409	19.2967	4.3551	0.9313	0.1484	0.1077
Std	0.3354	5.1192	0.5742	0.0193	0.0199	0.0141

\*\*\* 1 - LGBMRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						

0	3.5398	29.0514	5.3899	0.9121	0.1635	0.1275
...						
9	2.9727	19.1409	4.3750	0.9252	0.1640	0.1160
Mean	3.1203	21.8118	4.6207	0.9222	0.1522	0.1100
Std	0.3422	6.4237	0.6787	0.0236	0.0230	0.0142

\*\*\* 2 - ExtraTreesRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						
0	5.3227	53.8647	7.3393	0.8370	0.2351	0.2000
1...						
9	5.0418	38.2233	6.1825	0.8506	0.2439	0.2172
Mean	4.9365	41.6118	6.4389	0.8523	0.2119	0.1816
Std	0.2409	5.1511	0.3898	0.0196	0.0268	0.0253

\*\*\* 3 - RandomForestRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						
0	4.6211	46.6803	6.8323	0.8588	0.2290	0.1846
...						
9	4.6621	32.9120	5.7369	0.8714	0.2293	0.2000
Mean	4.5181	35.5683	5.9522	0.8745	0.2030	0.1707
Std	0.1097	4.5604	0.3733	0.0121	0.0318	0.0272

\*\*\* 4 - GradientBoostingRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						
0	3.3277	25.1146	5.0114	0.9240	0.1740	0.1271
...						
9	2.9214	19.6030	4.4275	0.9234	0.1694	0.1215
Mean	3.1014	20.4411	4.4966	0.9272	0.1542	0.1114
Std	0.2730	4.1364	0.4706	0.0171	0.0203	0.0145

\*\*\* 5 - MLPRegressor \*\*\*

Fitting 10 folds for each of 24 candidates, totalling 240 fits

	MAE	MSE	RMSE	R2	RMSLE	MAPE
Fold						
0	5.6160	56.9794	7.5485	0.8276	0.2771	0.1929
...						
9	5.1128	37.7620	6.1451	0.8524	0.2474	0.2167
Mean	5.1949	46.8561	6.8208	0.8327	0.2217	0.1769
Std	0.4478	7.8475	0.5772	0.0343	0.0291	0.0239

Transformation Pipeline and Model Successfully Saved

\*\*\* Tuning and finalizing: 165.12 seconds

\*\*\* Total time: 183.60 seconds

(Yes, depending on the versions of the libraries that you have installed, there may be some warnings from this process. I'm not showing those here.)

When I run this, I end up with a pickle file for the top six models:

- `LGBMRegressor.pkl`
- `CatBoostRegressor.pkl`
- `MLPRegressor.pkl`
- `ExtraTreesRegressor.pkl`
- `RandomForestRegressor.pkl`
- `GradientBoostingRegressor.pkl`

## 1.2 Testing

You will write a program called `concrete_test.py` that will scan the current directory for `.pkl` files. It will use `pycaret` to load those in one at a time.

Each model will be tested on `concrete_test.py`. The program will print the time that inference required and the  $R^2$  value.

Run the program and save the output to `test.txt`

My `test.txt` looks like this:

```
GradientBoostingRegressor:
  Inference: 0.0095 seconds
  R2 on test data = 0.9110
RandomForestRegressor:
  Inference: 0.0319 seconds
  R2 on test data = 0.8815
AdaBoostRegressor:
  Inference: 0.0147 seconds
  R2 on test data = 0.7239
ExtraTreesRegressor:
  Inference: 0.0170 seconds
  R2 on test data = 0.9007
MLPRegressor:
  Inference: 0.0052 seconds
  R2 on test data = 0.7861
CatBoostRegressor:
  Inference: 0.0030 seconds
  R2 on test data = 0.9079
LGBMRegressor:
```

Inference: 0.0071 seconds  
R2 on test data = 0.9101

Which would you use if accuracy was most important? What if speed was also really important?

## 2 $X^2$ testing for independence between categorical variables

Some times we will look at two categorical variables and try to figure out if they are related. Does knowing that the mouse has a particular gene tell us anything about the probability that it will get cancer?

You are given a csv with the results of this sort of experiment called `mice.csv`. Write a program `check_mice.py` that does the analysis. Put the analysis into a LaTeX file. (`mice.tex`) Convert that to a PDF (`mice.pdf`). Include bot files in your zip file.

For example, you should start out with a contingency table: (I did these examples with different data.)

Gene	No Cancer	Has Cancer	
R	34	2	<b>36</b>
J	4	45	<b>49</b>
K	17	18	<b>35</b>
	<b>55</b>	<b>65</b>	<b>120</b>

Then show conditional proportions:

Gene	No Cancer	Has Cancer	
R	94.4%	5.6%	<b>30.0%</b>
J	8.2%	91.8%	<b>40.8%</b>
K	48.6%	51.4%	<b>29.2%</b>
	<b>45.8%</b>	<b>54.2%</b>	

Then show the expected counts if the gene and cancer were independent:

Gene	No Cancer	Has Cancer	
R	16.5	19.5	<b>36</b>
J	22.5	26.5	<b>49</b>
K	16.0	19.0	<b>35</b>
	<b>45.8%</b>	<b>54.2%</b>	

Use the two tables to find  $X^2$ :

$$X^2 = 62.379$$

Note the degrees of freedom. (It is 2.)

And do a p-test:

$$p = 2.853273173286652 \times 10^{-14}$$

And then give proclamation: "It seems very, very unlikely that we would have seen these numbers if the gene and cancer were independent."

### 3 Criteria for success

If your name is Fred Jones, you will turn in a zip file called `HW10_Jones_Fred.zip` of a directory called `HW10_Jones_Fred`. It will contain:

- `concrete_train.py`
- `concrete_test.py`
- `check_mice.py`
- `test.txt`
- `train.txt`
- `mice.tex`
- `mice.pdf`
- `train_concrete.csv`
- `test_concrete.csv`
- `mice.csv`

Be sure to format your python code with black before you submit it.

We would run your code like this:

```
cd HW10_Jones_Fred
python3 concrete_train.py
python3 concrete_test.py
python3 check_mice.py
```

Do this work by yourself. Stackoverflow is OK. A hint from another student is OK. Looking at another student's code is *not* OK.

The template files for the python programs have import statements. Do not use any frameworks not in those import statements.